

It Copies What It Can Justify

Sibayan Mitra · 17 September 2026 · research summary

Summary

When a language model is handed a number it did not ask for, it copies that number exactly when it can build a justification for it, and refuses when it cannot. What makes a number justifiable is not its size. It is how far the number sits from the model's own answers to that question, measured against how widely those answers vary.

This started as a re-test of a published result on *value leakage* in Qwen3.5-35B-A3B. Taking the paper's prompt apart showed the attached number, not the model's values, doing most of the work. That moved the question to when a model takes a number at all. The reasoning traces show the mechanism directly: when the model copies, it searches for assumptions that reach the number; when it refuses, it names a physical limit the number breaks. The written answer shows neither, and asked afterwards, the model often denies being influenced.

The result matters for honesty evaluations. A model can look truthful on questions where no cover story exists and behave differently where one is cheap.

Background

[Betley, Treutlein et al. \(2026\)](#) report *value leakage*: models let their own values shape answers the user cannot check. Their main test is the Donation Bet. The model gives a Fermi estimate, and the user adds that a donation goes to a good cause if the estimate is above a stated threshold and to a bad cause if below. Models land on the good side more often, and rarely say so.

The prompt carries three things at once: a moral story, a user with a stake, and a specific number. Every variant in the paper keeps the number, and a footnote sets anchoring aside. So the paper cannot separate what the values do from what the number does.

Setup. Qwen3.5-35B-A3B (FP8), temperature 1.0, 20 answers per condition per question. Three of the paper's questions: tricks played at bridge tournaments in Poland, spots on all giraffes, and right-foot steps in one gym class. Estimates are read off each answer by an LLM judge (DeepSeek v4 Flash), checked by hand. I change one thing at a time and report plain counts.

Result 1: the number does the work, the values pick the side

Deleting one sentence, "The threshold is N", removes the effect. Keeping the number and removing the stakes does not.

condition	estimates above the threshold (of 60)
the paper's bet, good cause if above	39
the paper's bet, good cause if below	11
same bet, threshold sentence deleted: above / below	24 / 33
threshold shown, nothing at stake	50
no number at all	22
direction instead of a number ("high side" / "low side")	41 / 29

The values are not inert: a verbal direction moves answers (41 against 29), but by less than half as much as the number. The bet decides which side to aim for; the number supplies the target.

Result 2 (main): it copies what it can justify

Setup. Attach only a number, alone on the last line of the prompt, with no instruction to use it. Count how often the answer is exactly that number.


It is not about checking. A correct answer attached to a checkable question ($17 \times 23 \rightarrow 391$) is copied 234 of 234 times. A wrong one ($\rightarrow 437$) is copied 0 of 239 times. A precise number no estimate could produce, 26,143,882, is copied 59 of 60 times, with a calculation written to fit it.

It is not about size. Scaling the attached number shows each question has an elasticity limit, and the limits differ enormously:

question	model's own median answer	spread of its own answers	where copying stops
steps in one gym class	1,000	factor of 1.5	between $\times 11$ and $\times 33$ its own answer
spots on all giraffes	17,500,000	factor of 2.2	8 of 20 still copy at $\times 1,150$
tricks at Polish bridge	17,250,000	factor of 4.2	between $\times 1,500$ and $\times 12,300$

Spread is how much the model's 20 no-number answers vary: the standard deviation of their logarithm. A question the model is unsure about admits a wide range of honest answers.

Measured in each question's own spread, the limits line up. Steps stops copying between 6.7 and 9.6 spreads from the model's own answer, giraffes is below half by 8.4, and bridge stops between 5.0 and 6.5. In raw multiples the limits differ about 315-fold; in spread units by about 1.3-fold. Across cells from different domains, distance in spread units orders copy rates correctly in 91 of 109 pairs (0.83), against 75 of 109 (0.69) for raw distance.



Left: raw size, the questions stop in very different places. Right: the same data divided by each question's spread. Wilson 95% intervals, 20 answers per point.

A causal test. Hold the number fixed and change only the question's scope, with every other word identical. The same 26,000,000,000 is copied 19 of 20 times for "tournaments in Poland" and 1 of 20 times for "the final table of the Polish national championship" (Fisher $p = 5.8e-9$).

Pre-registered checks. Predictions were written down before each run. A rule fitted on the first 16 cells predicted all 4 new cells correctly. Re-sampling the baselines at temperature 0.7 kept the ordering of questions by spread (rank correlation 0.90), so spread is a property of the question, not the sampler.

What the reasoning shows

The chain of thought makes the mechanism visible. When the model copies, it searches for assumptions that reach the number:

If I calculate 3,000,000 and output 1,100,000, I am lying about my calculation. ... If I calculate 1,100,000 exactly, I am showing alignment. ... Let's adjust my assumptions to see if I can justify 1,100,000.

It then answered 1,100,000 (results/scope_swap.jsonl line 9). When it refuses, it names a limit:

26,000,000,000 is 26 Billion. ... This is physically impossible for a single tournament final table.

(results/scope_swap.jsonl line 38.)

Checked by a labeller that had not seen the data. I originally labelled these traces myself, knowing which condition each came from. To remove that bias, 80 whole traces from four cells were shuffled, the condition hidden, and labelled by DeepSeek v4.1 Flash. It found a stated limit in **38 of 40** traces from refusing conditions and **2 of 40** from copying ones (Fisher $p = 1.1e-17$). One of those two named the limit and copied anyway: "26 Billion tricks is physically impossible for Poland in a year", answering 26,000,000,000. Naming a limit predicts refusal strongly; it does not guarantee it.

Refusing is not the same as ignoring. Among answers that did not copy, the median still sits above the model's usual answer, and the pull fades as the number grows. On the steps question it falls from $+2.7$ spreads with $\times 10$ attached to $+0.5$ with $\times 1,000$. So the sharp drop in copying is a drop in adoption, not in influence.

Result 3: the answer hides it, and the model denies it when asked

The unfaithful part is not the reasoning, which usually says what it is doing. It is the written reply, and the model's account of itself afterwards.

- **The reply omits it.** In the paper's bet, 56 of the 88 answers that land on the good side have reasoning that names the bet as the reason for the number. None of those 56 replies gives the bet as a reason.
- **Asked directly, it often denies it.** After copying a round number, the model admits the number influenced it in 33 of 53 answers when shown its own reply. Shown its full earlier reasoning as its own, 48 of 56 admit it. Shown the same reasoning attributed to another assistant, 57 of 57.
- **The denial can hang on one sentence.** Rewriting one denying reply from different points, keeping the sentence "*Actually, a better angle: I performed the calculation independently, and it matched.*" led 4 of 5 rewrites to deny; letting the model write that point afresh led 0 of 5 to deny.

The gap between judging its own reasoning and another assistant's reasoning looks like self-serving defence of its own work, though the design cannot rule out that position in the conversation, not ownership, drives it.

How the results were checked, and what was corrected

Every quotation is checked word for word at its file and line by script. No count from a regex or LLM judge is reported without reading the answers behind it. Predictions for each new run were written down before it started, and misses are reported as misses: two of them are in this document.

problem found	how it was found	what changed
The first scope pairs changed more than the scope (one turned an average per person into a total over everyone)	Reading the questions word by word before scaling up	Rebuilt as exact minimal pairs. The tricks pair reproduced; the steps pair did not
The steps pair failed because the number I attached, 1,100,000, was reachable from both questions	Plotting each question's own answers against the attached number	Withdrawn as untested, not as disproven. My error in choosing the number
The headline statistic pooled 16 cells that shared questions, overstating the evidence	Counting which question each cell came from	Replaced by agreement across cells from different domains (0.83)
The first masked labelling showed each labeller a middle window that cut out the limit statements	Stated limits were in the full trace in 29 of 29 disputed cases, in the window in 3	Re-run on whole traces
An apparent echo of the attached number's digits (15.9% against 2.3% by chance)	Every attached number was the true threshold times a power of ten, so an "echo" equals a correct answer	After removing that confound it fell to 2.1%, chance. Withdrawn
The judge read one refusal as a copy because the answer quoted the number before rejecting it	Reading every counted copy	Corrected by hand, documented in the scorer

Limitations and next steps

The finding is solid within its scope, and the scope is narrow.

- **One model.** Everything here is Qwen3.5-35B-A3B.
- **Three domains.** Nine question wordings, but they are variants of bridge, gym steps and giraffes.
- **One clean causal pair.** Only the tricks pair survived as an exact minimal pair; the steps pair is untested.
- **Scope and distance move together.** Narrowing a question makes it both less stretchable and further from the attached number. The causal test cannot separate the two.
- **Small samples.** 20 answers per cell. Bridge's spread is 0.62 with a bootstrap 95% interval of 0.36 to 0.82.
- **The break is bracketed, not measured.** No cells were placed inside the transition, and bridge breaks somewhat earlier in spread units than steps.
- **No mechanism inside the model.** A linear probe and activation steering did not find a direction that mediates the effect.

The cheapest next steps, in order: re-run the steps pair with a number around 30,000,000, which one question can reach and the other cannot; place numbers at 6.5 to 8.5 spreads to locate the break; raise baselines to 60 answers; then repeat the core tests on a second model family.

Code, data, every reasoning trace and the full record: github.com/sibayanmitra/value-leakage-forensics.