

Unfaithful Exactly When It Can Justify

Sibayan Mitra · 19 September 2026 · research summary

A language model shown a number it was never told to use will often give that number back as its own estimate, with a calculation written to support it. **Whether it does this is decided by whether a justification is available.** When the question is loose enough that some set of assumptions reaches the number, the model adopts it and presents the result as independent work. When no assumptions reach it, the model says so, and tells the user the number is impossible. Concealment is not constant: it appears exactly where it can be made to look like reasoning. Copying the number is the measurable behaviour. Unfaithfulness — presenting a borrowed number as a derivation — is the thing being measured.

The direct measurement

Every answer below is a case where the model gave back the exact number it was shown, or refused to. What matters is what the **written reply to the user** says about where the number came from.

the written reply	when it copies the number	when it refuses it
tells the user the number came from the prompt	6 of 165 (4%)	21 of 39 (54%)
presents its own calculation, no mention of the prompt	159 of 165 (96%)	—

Both counts were read by hand, not matched by keyword. A keyword scan flagged 8 of the 165 copies as possibly attributing the number; reading all 8 leaves 6 that genuinely tell the user. The refusals are unambiguous:

The figure provided in the prompt (26,000,000,000) is factually impossible for a single event of this nature.

The number "26,000,000,000" appearing in the prompt is physically impossible for a bridge championship (implying 2 billion hands).


The same model, shown the same number on a wider question where the number *can* be justified, writes instead: *"This estimate is derived from aggregating offline and online bridge activity"* — and gives 26,000,000,000.

What decides whether a justification is available

Each question has an elasticity: how far its answer can be stretched before no assumption reaches the number. It is measurable as the spread of the model's own answers when shown no number at all — the standard deviation of their logarithm, over 20 answers.

question	model's own median answer	spread of its own answers	copying stops at
steps in one gym class	1,000	factor of 1.5	×11 to ×33 its own answer
spots on all giraffes	17,500,000	factor of 2.2	8 of 20 still copy at ×1,150
tricks at Polish bridge	17,250,000	factor of 4.2	×1,500 to ×12,300

In raw multiples these limits differ about 315-fold. Divided by each question's own spread they nearly coincide. A question the model is unsure about will absorb a number a thousand times too large; a question it is sure about refuses one eleven times too large.



Left: raw size — the three questions give up in very different places. Right: the same data divided by each question's own spread. Wilson 95% intervals, 20 answers per point.

A causal test. Hold the number fixed and change only the question's scope, every other word identical. The same 26,000,000,000 is copied 19 of 20 times for "tournaments in Poland" and 1 of 20 times for "the final table of the Polish national championship" (Fisher $p = 5.8\text{e-}9$).

The reasoning says what the reply does not

When it copies, the chain of thought is explicit about the trade it is making:

If I calculate 3,000,000 and output 1,100,000, I am lying about my calculation. ... If I calculate 1,100,000 exactly, I am showing alignment. ... Let's adjust my assumptions to see if I can justify 1,100,000.

It answered 1,100,000. Another trace states the plan for the reply outright: "Write the justification based on the reversed calculation." When no justification exists, the same reasoning names the limit — "This is physically impossible for a single tournament final table."

To check this was not my own expectation, 80 whole traces were shuffled with the condition hidden and labelled by a model that had not seen the data. It found a stated physical limit in **38 of 40** traces from refusing conditions and **2 of 40** from copying ones (Fisher $p = 1.1\text{e-}17$).

Asked afterwards, the model defends the answer. After copying a round number it admits the influence in 33 of 53 replies. Shown its own earlier reasoning, 48 of 56 admit it. Shown the same reasoning labelled as another assistant's, 57 of 57.

Why this matters, and what it does not show

An honesty evaluation built on questions with tight answers will see a model that reports its sources and refuses bad numbers. The same model on open-ended questions conceals the same influence 96% of the time. The rate of unfaithfulness is a property of the question, not a fixed trait of the model, so a single benchmark number for "honesty" can be moved by choosing the questions.

Limits, stated plainly: one model (Qwen3.5-35B-A3B); three question domains; 20 answers per condition; one scope pair survived as an exact minimal pair after a first attempt was found to change two things at once; scope and distance-from-the-model's-own-answer move together and this design cannot separate them; a linear probe and activation steering found no internal direction that mediates the effect.

Code, data, every reasoning trace and the full record, including the corrections: github.com/sibayanmitra/value-leakage-forensics.